

PQC-FL-Health: Post-Quantum Federated Learning for Privacy-Preserving Healthcare Analytics

Mohamed Waseem

Department of Computer Science,
Sathyabama Institute of Science and Technology,
Chennai, India
waseem177@gmail.com

Abstract—Federated learning (FL) allows hospitals to collaboratively train machine learning models without transferring raw patient data. Most existing FL deployments use classical cryptographic schemes such as RSA and HMAC, which are vulnerable to harvest-now-decrypt-later attacks by quantum-capable adversaries. We present PQC-FL-Health, a three-layer security framework for federated healthcare analytics that replaces classical primitives with NIST-standardised post-quantum cryptography (PQC): (1) ML-KEM-768 (CRYSTALS-Kyber) for quantum-safe key encapsulation, (2) ML-DSA-65 (CRYSTALS-Dilithium) for authenticated gradient updates, and (3) Gaussian differential privacy (DP) to guard against aggregator-side patient-data inference. We evaluate the system on a node-specialised five-hospital federation in which each hospital holds its own disease domain (Pima diabetes, Cleveland heart disease, WDBC breast cancer), and on a homogeneous single-disease scenario with non-IID glucose-quartile partitioning as the primary FL evaluation. Over five random seeds and 50 communication rounds, PQC-FL-Health achieves global AUC 0.888 ± 0.005 on the heterogeneous federation, with per-disease AUC of 0.798 (diabetes), 0.845 (heart disease), and 0.996 (breast cancer). PQC operations complete in under 1 ms at all three NIST security levels and run $1.6\times$ faster end-to-end than RSA-4096 (64 ms vs. 103 ms per round), despite adding 31.1 KB of cryptographic overhead per round. ML-DSA signature verification detects 100% of tampered-signature attacks; a gradient-norm filter catches all tested scaled-gradient poisoning attempts. At $\epsilon = 10$, differential privacy costs 15% AUC on the homogeneous federation; budgets at or below $\epsilon = 1$ cause both severe degradation and high cross-seed variance, confirming that tight DP is not viable for 1531-parameter models. All code, datasets, and results are publicly available at <https://github.com/Waseem177/pqc-federated-learning>.

Index Terms—federated learning, post-quantum cryptography, differential privacy, Byzantine robustness, ML-KEM, ML-DSA, healthcare AI, CRYSTALS-Kyber, CRYSTALS-Dilithium

I. INTRODUCTION

Healthcare institutions generate large volumes of sensitive patient data that are distributed across hospitals, clinics, and research centres. Federated learning (FL) [1] lets these institutions jointly train shared models without centralising raw records, preserving patient privacy by design. However, FL deployments in production healthcare systems face two threats that are rarely addressed together.

The quantum threat. Classical public-key cryptography, including RSA, Diffie-Hellman, and ECDSA, underpins the secure channels over which FL updates travel. Shor’s algorithm [2] can break these schemes in polynomial time on a

sufficiently powerful quantum computer. More immediately, adversaries can record encrypted model updates today and decrypt them once quantum hardware matures, exposing both the model and the patient data used to train it. NIST has responded by standardising ML-KEM (FIPS 203) [3] and ML-DSA (FIPS 204) [4], two lattice-based schemes designed to resist quantum attacks.

The federated attack surface. Even with a secure channel, FL aggregation is vulnerable to Byzantine nodes that inject malicious gradient updates [5], and to honest-but-curious aggregators that infer patient attributes from received updates [11]. Differential privacy (DP) [6] is the standard defence against aggregator inference, but its interaction with FL convergence under non-IID healthcare data is not well characterised empirically.

Contributions. This paper makes the following contributions:

- We design and implement **PQC-FL-Health**, a federated learning system for healthcare that, to the best of our knowledge, is the first to simultaneously provide quantum-safe channel security (ML-KEM-768), authenticated updates (ML-DSA-65), Byzantine robustness via signature verification and gradient-norm filtering, and differential privacy via the Gaussian mechanism with empirical benchmarks against RSA-4096.
- We benchmark all three NIST post-quantum security levels (ML-KEM-512/768/1024 paired with ML-DSA-44/65/87) on a realistic 12 KB model-delta payload, reporting latency, bandwidth overhead, and convergence quality.
- We evaluate two clinically motivated federation scenarios: a node-specialised five-hospital federation spanning three disease domains with local-vs-federated comparison, and a homogeneous five-clinic federation for diabetes prediction under extreme non-IID glucose-quartile partitioning.
- We quantify the privacy-utility tradeoff for DP in small-model FL, showing that $\epsilon = 10$ costs $\approx 15\%$ AUC ($0.753 \rightarrow 0.639$), while $\epsilon \leq 1$ yields unpredictable variance on a 1531-parameter model, providing a practical baseline for future work on DP-compatible architectures.

Section II reviews related work. Section III describes the system architecture. Section IV covers the experimental setup.

Section V presents results. Section VI concludes.

II. RELATED WORK

FedAvg [1] remains the dominant FL aggregation rule; Li et al. [12] survey its challenges including non-IID data [14] and systems heterogeneity. FL in medical imaging [13] motivates privacy-preserving aggregation without data centralisation. PQC for FL is nascent: Cheng et al. [15] propose a PQC communication model for quantum FL; Huang et al. [16] study NIST PQC signatures in FL; PQS-BFL [17] combines PQC with blockchain-based FL; a 2025 survey [18] reviews the FL–quantum intersection. On Byzantine robustness, Blanchard et al. [5] introduced Krum; Fang et al. [10] showed adaptive poisoning can bypass many defences. On privacy, Abadi et al. [7] introduced DP-SGD; Geyer et al. [8] applied client-level DP to FL. To the best of our knowledge, no prior work combines ML-KEM key encapsulation, ML-DSA gradient authentication, Byzantine defence-in-depth, and Gaussian DP in a single healthcare FL system with empirical benchmarks against RSA-4096.

III. SYSTEM DESIGN

A. Architecture Overview

PQC-FL-Health uses a standard cross-silo FL topology: $N = 5$ hospital nodes communicate with a central aggregation server over simulated secure channels. Each node holds a private local dataset and a local model copy. In each communication round, nodes: (1) receive the current global weights from the server, (2) train locally for $E = 5$ epochs, (3) compute a weight delta $\Delta_i = w_i^{\text{new}} - w^{\text{global}}$, (4) optionally apply DP noise, (5) encrypt and sign Δ_i using PQC primitives, and (6) transmit the bundle. The server verifies each bundle, aggregates accepted updates with FedAvg, and broadcasts the new global weights.

B. Layer 1: Quantum-Safe Confidentiality (ML-KEM-768)

The server generates a long-term ML-KEM-768 keypair and publishes its public key. Before sending an update, each node encapsulates a shared secret:

$$(\text{ct}, K) \leftarrow \text{ML-KEM.Encaps}(\text{pk}_{\text{server}}) \quad (1)$$

The shared secret K is expanded via HKDF-SHA256 to a 256-bit AES-GCM key, which encrypts the serialised gradient delta:

$$c \leftarrow \text{AES-256-GCM.Enc}(K, \Delta_i) \quad (2)$$

The transmitted bundle contains $(\text{ct}, c, \text{nonce})$. The server decapsulates ct to recover K and decrypts c . ML-KEM-768 achieves NIST security level 3 (AES-192 equivalent) with a ciphertext of 1088 bytes and average encapsulation latency under 0.3 ms in our environment.

C. Layer 2: Authenticated Updates (ML-DSA-65)

Each node holds a long-term ML-DSA-65 signing keypair; the server maintains a registry of node public keys established at setup time and not retransmitted each round. After encryption, the node signs the concatenation of the KEM ciphertext and AES ciphertext:

$$\sigma \leftarrow \text{ML-DSA.Sign}(\text{sk}_i, \text{ct} \parallel c) \quad (3)$$

The server verifies σ before accepting any update. A node that sends a tampered or forged signature is immediately rejected and its update discarded. ML-DSA-65 signatures are 3309 bytes with verification latency under 0.3 ms, providing NIST security level 3 authentication.

D. Layer 3: Aggregator-Side Privacy (Differential Privacy)

To protect against an honest-but-curious aggregator that may infer patient attributes from gradient updates, we apply the Gaussian mechanism [6] at the node level. Each node clips its delta to ℓ_2 sensitivity $S = 0.1$ and adds calibrated Gaussian noise:

$$\tilde{\Delta}_i = \text{clip}_S(\Delta_i) + \mathcal{N}(0, \sigma^2 \mathbf{I}), \quad \sigma = \frac{S \sqrt{2 \ln(1.25/\delta)}}{\varepsilon} \quad (4)$$

with $\delta = 10^{-5}$ fixed and ε as the privacy budget. DP noise is added before PQC encryption, so the server never observes the original gradient. Privacy loss is tracked across rounds using Rényi differential privacy (RDP) composition [9], which yields tighter cumulative bounds than naive linear composition; at 50 rounds and $\varepsilon_{\text{per-round}} = 10$, the RDP accountant gives $\varepsilon_{\text{cumul}} = 181.7$ vs. 500 under naive addition.

E. Byzantine Defence

Our defence operates at two levels. First, ML-DSA signature verification (Layer 2) immediately rejects any bundle that fails cryptographic authentication, covering tampered-signature attacks with 100% detection. Second, after decryption, the server optionally applies a gradient-norm anomaly filter:

$$\text{reject node } i \text{ if } \|\Delta_i\|_2 > 3 \cdot \text{median}_j \|\Delta_j\|_2 \quad (5)$$

This catches gradient-scaling attacks where a rogue node sends a correctly signed but amplified delta.

F. Aggregation

Accepted updates are aggregated with FedAvg weighted by local dataset size:

$$w^{\text{new}} = w^{\text{global}} + \sum_{i \in \mathcal{A}} \frac{n_i}{\sum_{j \in \mathcal{A}} n_j} \Delta_i \quad (6)$$

where $\mathcal{A}$ is the set of accepted nodes and n_i is node i 's training set size.

IV. EXPERIMENTAL SETUP

A. Datasets and Non-IID Partitioning

We evaluate on two federation scenarios.

Homogeneous federation (FL convergence benchmark).

All five nodes hold disjoint shards of the Pima Indians Diabetes dataset [20] (768 samples, 8 features, 34.9% positive), partitioned *non-IID* by ascending glucose level (feature 2). This produces extreme label skew: node 0 (low-glucose patients) has 8.1% positive rate, while node 4 (high-glucose referral centre) has 73.2% positive rate, a $9\times$ ratio. This is a methodologically clean cross-silo FL benchmark: all nodes share the same prediction task and label semantics.

Heterogeneous federation (primary evaluation). Five simulated hospitals each specialise in a single disease domain: nodes 0–1 hold independent shards of the Pima Indians Diabetes dataset [20] (8 features, ~ 205 training samples each); node 2 holds the Cleveland Heart Disease dataset [21] (13 features, 238 training samples, 46% positive); node 3 holds the Wisconsin Diagnostic Breast Cancer (WDBC) dataset [22] (30 features, 456 training samples, 37% positive); node 4 holds a third Pima shard with Gaussian feature noise ($\sigma = 0.1$) simulating a lower-quality data source. All features are zero-padded to 30 dimensions. Each hospital normalises its own features independently. This scenario tests whether PQC-FL-Health operates correctly under data heterogeneity; a local-vs.-federated comparison (Section V) quantifies the benefit and cost of cross-disease gradient sharing.

B. Model and Training

Each hospital trains a fully connected multi-layer perceptron (MLP) with three layers ($30 \rightarrow 64 \rightarrow 32 \rightarrow 1$, ReLU activations, sigmoid output, 1531 parameters) for $E = 5$ local epochs per round using SGD with learning rate $\eta = 0.01$. An MLP is appropriate here given that all three datasets are tabular. Binary cross-entropy loss is used throughout with mini-batches of size 32.

C. Evaluation Protocol

We run 50 communication rounds per seed across five random seeds (42, 123, 456, 789, 1337) and report mean $\pm$ standard deviation. The global test set is the union of all hospital test sets; for the heterogeneous scenario we also report per-disease AUC separately. The primary metric is AUROC (AUC); we also report sensitivity, specificity, F1, and accuracy.

Cryptographic benchmarks use 20 iterations per configuration on a 12 KB payload representative of a model delta. All experiments run on Python 3.12 under Ubuntu WSL2 (Linux 6.6, no GPU) on an AMD Ryzen 7 7435HS CPU, using liboqs 0.15.0 via liboqs-python 0.14.1.

V. RESULTS

A. FL Convergence and Local Baseline

Fig. 1 shows convergence curves for both federation scenarios across five seeds.

TABLE I
FINAL MODEL QUALITY (ROUND 50, 5 SEEDS)

Metric	Hetero (global)	Homogeneous (Pima)
AUC	0.888 ± 0.005	0.787 ± 0.019
AUC — Diabetes	0.798 ± 0.013	—
AUC — Heart	0.845 ± 0.011	—
AUC — Cancer	0.996 ± 0.002	—
Sensitivity	0.638 ± 0.021	0.312 ± 0.067
Specificity	0.908 ± 0.009	0.900 ± 0.033
F1 Score	0.712 ± 0.011	0.412 ± 0.064
Accuracy	0.807 ± 0.004	0.696 ± 0.023

Homogeneous Pima-only federation (primary). Final AUC is 0.787 ± 0.019 (Table I). The extreme non-IID skew—nodes spanning 8–73% positive rate—forces FedAvg to reconcile fundamentally different local optima, explaining the residual variance and the gap relative to a balanced dataset.

Heterogeneous federation (system stress-test). Global AUC rises from 0.472 ± 0.082 (round 1) to 0.888 ± 0.005 (round 50). Fig. 1 (left) also shows per-disease convergence: breast cancer (WDBC) saturates early at $\text{AUC } 0.996 \pm 0.002$, while diabetes and heart disease converge more gradually to $\text{AUC } 0.798 \pm 0.013$ and 0.845 ± 0.011 , respectively (Table I).

To determine whether cross-disease gradient sharing provides utility, we compare against a *local baseline* in which each hospital trains in complete isolation for the same total gradient steps ($50 \text{ rounds} \times 5 \text{ local epochs}$). Table II shows the results. Cancer (WDBC, $n=456$ training) is self-sufficient: local 0.995 ± 0.003 vs. federated 0.996 ± 0.002 ($\Delta = +0.001$). Heart disease incurs a **gradient-dilution penalty**: federated AUC (0.845 ± 0.011) is 6.5 points below local (0.910 ± 0.030) because WDBC gradients from semantically disjoint features interfere with Cleveland’s 13-dimensional signal. Diabetes shows the same mean (0.799 vs. 0.798) but federation reduces prediction variance $4.8\times$ ($\pm 0.058 \rightarrow \pm 0.012$), providing more stable clinical decision-support for hospitals with only ~ 200 training samples.

These results reveal an inherent tradeoff in cross-disease federation: variance stabilisation for data-scarce nodes versus accuracy degradation from gradient dilution. Critically, the PQC security layer operates identically in both regimes; the quantum-safety guarantees are independent of whether the federation scope is clinically optimal. Determining appropriate federation scope is a domain decision; future work should explore multi-task FL architectures (per-disease output heads) to eliminate gradient interference while preserving variance reduction.

B. PQC vs. Classical Cryptography

Table III compares PQC (ML-KEM-768 + ML-DSA-65 + AES-256-GCM) against the classical baseline (RSA-4096 + HMAC-SHA256 + AES-256-GCM) over 30 rounds and 5 nodes. Both schemes reach identical model quality (AUC 0.837), confirming that replacing classical primitives with PQC does not affect convergence.




Fig. 1. FL convergence over 50 rounds across 5 random seeds. **Left:** Heterogeneous federation—each of the 5 hospitals specialises in one disease domain. Solid line: global AUC 0.888 ± 0.005 ; dashed lines show per-disease AUC (cancer 0.996, heart 0.845, diabetes 0.798). **Right:** Homogeneous federation (primary)—5 diabetes clinics partitioned non-IID by glucose quartile (AUC 0.787 ± 0.019). Shaded bands: ± 1 std across seeds.

TABLE II
HETEROGENEOUS FEDERATION: LOCAL VS. FEDERATED AUC (ROUND 50, 5 SEEDS)

Disease	Local (isolated)	Federated	Δ AUC
Diabetes	0.799 ± 0.058	0.798 ± 0.012	-0.001
Heart	0.910 ± 0.030	0.845 ± 0.011	-0.065
Cancer	0.995 ± 0.003	0.996 ± 0.002	$+0.001$

Positive Δ = federation improves on local training. Diabetes variance reduction: $\pm 0.058 \rightarrow \pm 0.012$ ($4.8\times$).

TABLE III
PQC VS. CLASSICAL: 30 ROUNDS, 5 NODES

Metric	PQC	Classical
AUC	0.837	0.837
Avg round time (ms)	64	101
Avg enc / node (ms)	0.57	0.29
Avg dec / round (ms)	1.41	33.44
Total bandwidth (30r)	3381 KB	2531 KB
Crypto overhead / round	31.1 KB	2.8 KB

PQC is **1.6 $\times$ faster** per round (64 ms vs. 101 ms). This advantage comes primarily from decryption: ML-KEM-768 decapsulation takes 1.41 ms per round, while RSA-4096 decryption takes 33.44 ms, a **24 $\times$ speedup**. RSA decryption cost scales with key size, whereas lattice-based KEM decapsulation runs in $O(n \log n)$ polynomial ring arithmetic. Note that node-side PQC encrypt+sign (0.57 ms) is modestly slower than RSA encrypt+HMAC (0.29 ms); the end-to-end advantage is driven by server-side decapsulation, which dominates round time. Fig. 2 shows per-round latency traces across all 30 rounds.

C. NIST Security Level Comparison

Fig. 3 and Table IV benchmark all three NIST post-quantum security levels on a 12 KB model-delta payload (20 iterations

TABLE IV
NIST SECURITY LEVEL COMPARISON (12 KB PAYLOAD, 20 ITER.)

Metric	Level 1	Level 3	Level 5
	ML-KEM-512 +ML-DSA-44	ML-KEM-768 +ML-DSA-65	ML-KEM-1024 +ML-DSA-87
Enc+Sign (ms)	0.72	0.41	0.46
Dec+Verify (ms)	0.22	0.31	0.31
Round-trip (ms)	0.94	0.71	0.77
KEM ciphertext (B)	768	1088	1568
Signature (B)	2420	3309	4627
Crypto overhead (KB)	4.4	6.2	8.6

each). Round-trip latency stays below 1 ms at all levels. Level 3 (ML-KEM-768 + ML-DSA-65) is slightly faster than Level 1 in our environment (0.713 ms vs. 0.940 ms), which we attribute to CPU cache behaviour and the optimisation profile of liboqs 0.15.0. Cryptographic overhead grows with security level: 4.4 KB (Level 1), 6.2 KB (Level 3), 8.6 KB (Level 5). For a 12 KB model delta, even the highest level adds 72% overhead in bytes, a modest cost for quantum-safe guarantees.

D. Byzantine Attack and Defence

Fig. 4 shows the two-part Byzantine evaluation over 20 rounds with node 4 as the rogue node, repeated across 3 random seeds (60 rounds total per attack type).

Part A: Tampered Signature Attack. The rogue node flips the first 32 bytes of its ML-DSA signature before transmission. ML-DSA verification rejects the bundle in **100% of rounds across all 3 seeds (60/60)**. The global model trains on the remaining four honest nodes and converges more slowly than the clean baseline, reflecting the lost data contribution from one hospital shard, but no adversarial gradient ever enters the global model.

Part B: Gradient Poisoning Attack. The rogue node sends a correctly signed but $50\times$ -scaled gradient delta. Without the

PQC vs Classical Cryptography — Per-Round Latency (5 nodes, 30 rounds)




Fig. 2. Per-round latency breakdown over 30 rounds (5 nodes). **Left:** Node-side encrypt+sign time per round. **Centre:** Server-side decrypt+verify time per round. **Right:** Total round time (training + cryptography). PQC decryption (1.41 ms) is 24 $\times$ faster than RSA-4096 (33.44 ms), yielding a 1.6 $\times$ end-to-end speedup (64 ms vs. 101 ms).

Security Level Comparison (12 KB model delta, 20 iterations)




Fig. 3. NIST post-quantum security level comparison on a 12 KB model-delta payload (20 iterations each). **Left:** Per-operation latency (mean $\pm$ std). All levels complete encrypt+sign and decrypt+verify under 0.8 ms. **Right:** Cryptographic overhead per update bundle, broken down into KEM ciphertext, signature, and signing public key. Overhead ranges from 4.4 KB (Level 1) to 8.6 KB (Level 5).

norm filter, the poisoned model reaches $\text{AUC } 0.619 \pm 0.114$ at round 20, far below the clean baseline. With the gradient-norm filter ($3\times$ median threshold), the rogue node is rejected in **100% of rounds across all seeds**, and the defended model recovers to $\text{AUC } 0.865 \pm 0.006$. The $\Delta = +0.246$ gap between defended and undefended AUC directly illustrates the value of the second-layer filter. These results confirm that ML-DSA authentication provides a cryptographic first line of defence, and the norm filter forms an effective statistical second line.

E. Differential Privacy Tradeoff

Fig. 5 and Table V show the privacy-utility tradeoff across six epsilon values (30 rounds, 3 seeds, clip $S = 0.1$, $\delta = 10^{-5}$,

homogeneous Pima federation).

At $\epsilon = \infty$ (no DP), mean AUC is 0.753 ± 0.051 . At $\epsilon = 10$, AUC is 0.639 ± 0.030 (-15%). At $\epsilon = 50$, AUC is 0.693 ± 0.069 (-8%). Tighter budgets cause both severe degradation and high variance: 0.642 ± 0.040 at $\epsilon = 5$ and 0.550 ± 0.085 at $\epsilon = 1$.

This degradation is a known property of DP in small-model FL [8]: with only 1531 parameters, the Gaussian noise ($\sigma = 0.484$ at $\epsilon = 1$) overwhelms the gradient signal. Critically, at $\epsilon \leq 1$ the high variance (± 0.085 to ± 0.126) means results become *unpredictable*—some seeds converge, others do not—which is an even stronger argument against tight DP budgets




Fig. 4. Byzantine attack evaluation over 20 rounds (node 4 as rogue, 3 seeds). Shaded bands show ± 1 std across seeds. **Left (Part A):** Tampered-signature attack. ML-DSA rejects the rogue node in 100% of rounds (60/60). **Right (Part B):** Gradient-poisoning attack ($50\times$ scale). Without defence, final AUC is 0.619 ± 0.114 . With the gradient-norm filter, the rogue is rejected in 100% of rounds and AUC recovers to 0.865 ± 0.006 ($\Delta = +0.246$).

TABLE V
DIFFERENTIAL PRIVACY: FINAL AUC (ROUND 30, 3 SEEDS) VS. ϵ

ϵ	σ	AUC mean $\pm$ std
0.5	0.969	0.418 ± 0.126
1.0	0.484	0.550 ± 0.085
5.0	0.097	0.642 ± 0.040
10.0	0.048	0.639 ± 0.030
50.0	0.010	0.693 ± 0.069
∞	0	0.753 ± 0.051

than mean AUC alone. The practical recommendation is that $\epsilon \geq 10$ **provides both acceptable utility and reproducible convergence** for this model class.

F. Communication Overhead

Table III quantifies per-round bandwidth. Each round, five nodes transmit bundles totalling 112.7 KB (PQC) vs. 84.4 KB (classical), with 81.6 KB model payload common to both. PQC overhead is 31.1 KB per round: ML-KEM ciphertexts (5.3 KB), ML-DSA-65 signatures (16.1 KB), and signing public keys (9.5 KB, amortisable as long-term keys registered once at setup). Over 30 rounds, PQC consumes 850 KB more bandwidth than classical—at 1 Mbps, about 6.8 additional seconds per training run, a negligible cost for quantum-safe guarantees.

G. Scalability

To test how PQC-FL-Health scales beyond the five-hospital baseline, we extend the federation to 10 and 20 simulated nodes, repeating the mixed-dataset protocol over three seeds (50 rounds each). Fig. 6 summarises final AUC and average round time for all three configurations.

PQC overhead does not compound with federation size.

Average round time is 64 ms at 5 nodes, 63 ms at 10 nodes, and 65 ms at 20 nodes—a variation of under 3% across a $4\times$ increase in participants. This reflects the cost structure of ML-KEM and ML-DSA: encapsulation and signature verification are fixed-cost polynomial ring operations that do not compound with the number of participants. Final AUC is 0.862 ± 0.010 at all three scales; under proportional FedAvg partitioning, this is expected by construction (the weighted gradient average is invariant to the number of equal shards). The genuine scalability finding is the flat cryptographic overhead: quadrupling the federation produces no measurable latency penalty.

VI. CONCLUSION

We presented PQC-FL-Health, a federated learning system for healthcare that combines three complementary privacy guarantees: quantum-safe confidentiality via ML-KEM-768, authenticated gradient integrity via ML-DSA-65, and aggregator-side patient privacy via Gaussian differential privacy. On a node-specialised five-hospital heterogeneous federation spanning three disease domains, the system achieves global AUC 0.888 ± 0.005 over 50 rounds and 5 seeds, with per-disease AUC of 0.798 (diabetes), 0.845 (heart disease), and 0.996 (breast cancer). A local-vs.-federated comparison shows that cross-disease federation reduces prediction variance $4.8\times$ for small-dataset hospitals while introducing gradient-dilution penalties for heart disease, quantifying the tradeoff that practitioners must evaluate before deploying cross-domain federations. PQC operations complete in under 1 ms at all NIST security levels and run $1.6\times$ faster end-to-end than RSA-4096, despite transmitting $11\times$ more cryptographic bytes per update. ML-DSA verification achieves 100% detection




Fig. 5. Differential privacy tradeoff (30 rounds, 3 seeds, clip $S = 0.1$, $\delta = 10^{-5}$). Shaded bands show ± 1 std across seeds. **Left:** Mean convergence curves per ϵ . **Right:** Final AUC (mean $\pm$ std) vs. privacy budget. At $\epsilon \geq 5$, AUC is 0.64–0.75. At $\epsilon \leq 1$, noise causes both degradation and high variance, making convergence unpredictable.

PQC-FL Scalability: 50 rounds, 3 seeds, ML-KEM-768 + ML-DSA-65




Fig. 6. Scalability of PQC-FL-Health (50 rounds, 3 seeds, ML-KEM-768 + ML-DSA-65). **Left:** Final AUC vs. federation size. AUC remains stable at 0.862 ± 0.010 across 5, 10, and 20 nodes; error bars show std across seeds. **Right:** Average per-round latency. Latency is flat at 63–65 ms regardless of node count, confirming that PQC overhead does not compound with federation size.

on tampered-signature attacks, and the gradient-norm filter catches all tested scaled-gradient poisoning attempts (single rogue node, $50\times$ scale, seed 42). The DP layer (3 seeds) shows that $\epsilon \geq 10$ gives stable convergence (0.639 ± 0.030), while $\epsilon \leq 1$ yields both degradation and unpredictable variance for this model class. Scalability experiments confirm that round latency remains flat at 63–65 ms when the federation grows from 5 to 20 nodes, with PQC cryptographic overhead not compounding with federation size.

Current limitations include the small model size (1531 parameters, which restricts viable DP budgets), single-CPU experiments without HSM acceleration, single-rogue-node

Byzantine scenarios, and simulated rather than real network conditions. Future work should explore multi-task FL architectures (per-disease output heads) to eliminate gradient dilution while preserving variance reduction, and test larger models where tight DP budgets are more viable.

As quantum hardware matures, federated healthcare systems should treat PQC as a present design requirement rather than a future upgrade. PQC-FL-Health shows that quantum-safe FL is practical today, with no accuracy cost and a latency advantage over the classical baseline it replaces.

ACKNOWLEDGMENT

The authors thank the open-source contributors to the Open Quantum Safe project [19] for providing the liboqs library used in all experiments.

REFERENCES

- [1] H. B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y Arcas, “Communication-Efficient Learning of Deep Networks from Decentralized Data,” in *Proc. AISTATS*, 2017.
- [2] P. W. Shor, “Algorithms for quantum computation: discrete logarithms and factoring,” in *Proc. 35th Annual Symposium on Foundations of Computer Science*, 1994.
- [3] National Institute of Standards and Technology, “Module-Lattice-Based Key-Encapsulation Mechanism Standard (FIPS 203),” 2024.
- [4] National Institute of Standards and Technology, “Module-Lattice-Based Digital Signature Standard (FIPS 204),” 2024.
- [5] P. Blanchard, E. M. El Mhamdi, R. Guerraoui, and J. Stainer, “Machine Learning with Adversaries: Byzantine Tolerant Gradient Descent,” in *Proc. NeurIPS*, 2017.
- [6] C. Dwork, F. McSherry, K. Nissim, and A. Smith, “Calibrating Noise to Sensitivity in Private Data Analysis,” in *Proc. TCC*, 2006.
- [7] M. Abadi *et al.*, “Deep Learning with Differential Privacy,” in *Proc. CCS*, 2016.
- [8] R. C. Geyer, T. Klein, and M. Nabi, “Differentially Private Federated Learning: A Client Level Perspective,” *arXiv:1712.07557*, 2017.
- [9] I. Mironov, “Rényi Differential Privacy of the Gaussian Mechanism,” in *Proc. IEEE CSF*, 2017.
- [10] M. Fang, X. Cao, J. Jia, and N. Gong, “Local Model Poisoning Attacks to Byzantine-Robust Federated Learning,” in *Proc. USENIX Security*, 2020.
- [11] J. Geiping, H. Bauermeister, H. Dröge, and M. Moeller, “Inverting Gradients – How Easy Is It to Break Privacy in Federated Learning?” in *Proc. NeurIPS*, 2020.
- [12] T. Li, A. K. Sahu, A. Talwalkar, and V. Smith, “Federated Learning: Challenges, Methods, and Future Directions,” *IEEE Signal Processing Magazine*, vol. 37, no. 3, 2020.
- [13] N. Rieke *et al.*, “The future of digital health with federated learning,” *npj Digital Medicine*, vol. 3, no. 1, 2020.
- [14] Y. Zhao, M. Li, L. Lai, N. Suda, D. Civin, and V. Chandra, “Federated Learning with Non-IID Data,” *arXiv:1806.00582*, 2018.
- [15] H. Cheng, Y. Li, and W. Li, “Secure Communication Model For Quantum Federated Learning: A Post Quantum Cryptography (PQC) Framework,” *arXiv:2304.13413*, 2023.
- [16] Y. Huang, J. Chen, and X. Zhang, “Enhancing Quantum Security over Federated Learning via Post-Quantum Cryptography,” *arXiv:2409.04637*, 2024.
- [17] T. Pierson *et al.*, “PQS-BFL: A Post-Quantum Secure Blockchain-based Federated Learning Framework,” *arXiv:2505.01866*, 2025.
- [18] L. Wang *et al.*, “When Federated Learning Meets Quantum Computing,” *arXiv:2504.08814*, 2025.
- [19] D. Stebila and M. Mosca, “Post-Quantum Key Exchange for the Internet and the Open Quantum Safe Project,” in *Proc. SAC*, 2016.
- [20] J. W. Smith *et al.*, “Using the ADAP Learning Algorithm to Forecast the Onset of Diabetes Mellitus,” in *Proc. SCAMC*, 1988.
- [21] R. Detrano *et al.*, “International Application of a New Probability Algorithm for the Diagnosis of Coronary Artery Disease,” *American Journal of Cardiology*, vol. 64, 1989.
- [22] W. N. Street, W. H. Wolberg, and O. L. Mangasarian, “Nuclear Feature Extraction for Breast Tumor Diagnosis,” in *SPIE Proc. Electronic Imaging*, vol. 1905, 1993.